

Introduction to Hidden Markov Models


Markov Models

- Set of states: $\{s_1, s_2, \dots, s_N\}$
- Process moves from one state to another generating a sequence of states: $s_{i1}, s_{i2}, \dots, s_{ik}, \dots$
- Markov chain property: probability of each subsequent state depends only on what was the previous state:

$$P(s_{ik} | s_{i1}, s_{i2}, \dots, s_{ik-1}) = P(s_{ik} | s_{ik-1})$$

- To define Markov model, the following probabilities have to be specified: transition probabilities $a_{ij} = P(s_i | s_j)$ and initial probabilities $\pi_i = P(s_i)$

Example of Markov Model



- Two states: 'Rain' and 'Dry'.
- Transition probabilities: $P('Rain'|'Rain')=0.3$, $P('Dry'|'Rain')=0.7$, $P('Rain'|'Dry')=0.2$, $P('Dry'|'Dry')=0.8$
- Initial probabilities: say $P('Rain')=0.4$, $P('Dry')=0.6$.

Calculation of sequence probability

- By Markov chain property, probability of state sequence can be found by the formula:

$$\begin{aligned} P(s_{i1}, s_{i2}, \dots, s_{ik}) &= P(s_{i1}) P(s_{i2} | s_{i1}) P(s_{i3} | s_{i1}, s_{i2}) \dots P(s_{ik} | s_{i1}, s_{i2}, \dots, s_{ik-1}) \\ &= P(s_{i1}) P(s_{i2} | s_{i1}) P(s_{i3} | s_{i2}) \dots P(s_{ik} | s_{ik-1}) \\ &= P(s_{i1}) P(s_{i2} | s_{i1}) P(s_{i3} | s_{i2}) \dots P(s_{ik} | s_{ik-1}) \end{aligned}$$


- Suppose we want to calculate a probability of a sequence of states in our example, $\{ 'Dry', 'Dry', 'Rain', 'Rain' \}$.

$$\begin{aligned} P('Rain'|'Rain') P('Rain'|'Dry') P('Dry'|'Dry') P('Dry') &= \\ = 0.3 * 0.2 * 0.8 * 0.6 \end{aligned}$$

Hidden Markov models.

- Set of states: $\{s_1, s_2, \dots, s_N\}$
 - Process moves from one state to another generating a sequence of states: $s_{i1}, s_{i2}, \dots, s_{ik}, \dots$
 - Markov chain property: probability of each subsequent state depends only on what was the previous state:
- $$P(s_{ik} | s_{i1}, s_{i2}, \dots, s_{ik-1}) = P(s_{ik} | s_{ik-1})$$
- States are not visible, but each state randomly generates one of M observations (or visible states) $\{v_1, v_2, \dots, v_M\}$
 - To define hidden Markov model, the following probabilities have to be specified: matrix of transition probabilities $A=(a_{ij})$, $a_{ij} = P(s_i | s_j)$, matrix of observation probabilities $B=(b_i(v_m))$, $b_i(v_m) = P(v_m | s_i)$ and a vector of initial probabilities $\pi=(\pi_i)$, $\pi_i = P(s_i)$. Model is represented by $M=(A, B, \pi)$.

Example of Hidden Markov Model



Example of Hidden Markov Model

- Two states: 'Low' and 'High' atmospheric pressure.
- Two observations: 'Rain' and 'Dry'.
- Transition probabilities: $P('Low'|'Low')=0.3$, $P('High'|'Low')=0.7$, $P('Low'|'High')=0.2$, $P('High'|'High')=0.8$
- Observation probabilities: $P('Rain'|'Low')=0.6$, $P('Dry'|'Low')=0.4$, $P('Rain'|'High')=0.4$, $P('Dry'|'High')=0.6$
- Initial probabilities: say $P('Low')=0.4$, $P('High')=0.6$.

Calculation of observation sequence probability

- Suppose we want to calculate a probability of a sequence of observations in our example, $\{ 'Dry', 'Rain' \}$.
 - Consider all possible hidden state sequences:
- $$P('Dry', 'Rain') = P('Dry', 'Rain', 'Low', 'Low') + P('Dry', 'Rain', 'Low', 'High') + P('Dry', 'Rain', 'High', 'Low') + P('Dry', 'Rain', 'High', 'High')$$

where first term is:

$$\begin{aligned} P('Dry', 'Rain', 'Low', 'Low') &= \\ P('Dry', 'Rain' | 'Low', 'Low') P('Low', 'Low') &= \\ P('Dry'|'Low') P('Rain'|'Low') P('Low') P('Low'|'Low') &= \\ = 0.4 * 0.4 * 0.6 * 0.4 * 0.3 \end{aligned}$$

Main issues using HMMs:

Evaluation problem. Given the HMM $M=(A, B, \pi)$ and the observation sequence $O=O_1 O_2 \dots O_K$, calculate the probability that model M has generated sequence O .

Decoding problem. Given the HMM $M=(A, B, \pi)$ and the observation sequence $O=O_1 O_2 \dots O_K$, calculate the most likely sequence of hidden states S_i that produced this observation sequence O .

Learning problem. Given some training observation sequences $O=O_1 O_2 \dots O_K$ and general structure of HMM (numbers of hidden and visible states), determine HMM parameters $M=(A, B, \pi)$ that best fit training data.


$O=O_1 \dots O_K$ denotes a sequence of observations $O_k \in \{V_1, \dots, V_M\}$.

Word recognition example(1).

- Typed word recognition, assume all characters are separated.



- Character recognizer outputs probability of the image being particular character, $P(\text{image}|\text{character})$.



Word recognition example(2).


- Hidden states of HMM = characters.
- Observations = typed images of characters segmented from the image V_a . Note that there is an infinite number of observations
- Observation probabilities = character recognizer scores.

$$B = (b_i(v_a)) = (P(v_a | s_i))$$

- Transition probabilities will be defined differently in two subsequent models.

Word recognition example(3).


- If lexicon is given, we can construct separate HMM models for each lexicon word.



- Here recognition of word image is equivalent to the problem of evaluating few HMM models.
- This is an application of **Evaluation problem**.

Word recognition example(4).


- We can construct a single HMM for all words.
- Hidden states = all characters in the alphabet.
- Transition probabilities and initial probabilities are calculated from language model.
- Observations and observation probabilities are as before.



- Here we have to determine the best sequence of hidden states, the one that most likely produced word image.
- This is an application of **Decoding problem**.

Character recognition with HMM example.

- The structure of hidden states is chosen.




- Observations are feature vectors extracted from vertical slices.



- Probabilistic mapping from hidden state to feature vectors:
 1. use mixture of Gaussian models
 2. Quantize feature vector space.

Exercise: character recognition with HMM(1)

- The structure of hidden states:



- Observation = number of islands in the vertical slice.

- HMM for character 'A':

$$\text{Transition probabilities: } \{a_{ij}\} = \begin{pmatrix} .8 & .2 & 0 \\ 0 & .8 & .2 \\ 0 & 0 & 1 \end{pmatrix}$$

$$\text{Observation probabilities: } \{b_k\} = \begin{pmatrix} .9 & .1 & 0 \\ .1 & .8 & .1 \\ .9 & .1 & 0 \end{pmatrix}$$



- HMM for character 'B':

$$\text{Transition probabilities: } \{a_{ij}\} = \begin{pmatrix} .8 & .2 & 0 \\ 0 & .8 & .2 \\ 0 & 0 & 1 \end{pmatrix}$$

$$\text{Observation probabilities: } \{b_k\} = \begin{pmatrix} .9 & .1 & 0 \\ 0 & .2 & .8 \\ .6 & .4 & 0 \end{pmatrix}$$



Exercise: character recognition with HMM(2)

- Suppose that after character image segmentation the following sequence of island numbers in 4 slices was observed: { 1, 3, 2, 1 }

- What HMM is more likely to generate this observation sequence, HMM for 'A' or HMM for 'B'?

Exercise: character recognition with HMM(3)

Consider likelihood of generating given observation for each possible sequence of hidden states:

- HMM for character 'A':

Hidden state sequence	Transition probabilities	Observation probabilities
$S_1 \rightarrow S_1 \rightarrow S_2 \rightarrow S_3$	$.8 * .2 * .2$	$* .9 * 0 * .8 * .9 = 0$
$S_1 \rightarrow S_2 \rightarrow S_2 \rightarrow S_3$	$.2 * .8 * .2$	$* .9 * .1 * .8 * .9 = 0.0020736$
$S_1 \rightarrow S_2 \rightarrow S_3 \rightarrow S_3$	$.2 * .2 * 1$	$* .9 * .1 * .1 * .9 = 0.000324$
		Total = 0.0023976

- HMM for character 'B':

Hidden state sequence	Transition probabilities	Observation probabilities
$S_1 \rightarrow S_1 \rightarrow S_2 \rightarrow S_3$	$.8 * .2 * .2$	$* .9 * 0 * .2 * .6 = 0$
$S_1 \rightarrow S_2 \rightarrow S_2 \rightarrow S_3$	$.2 * .8 * .2$	$* .9 * .8 * .2 * .6 = 0.0027648$
$S_1 \rightarrow S_2 \rightarrow S_3 \rightarrow S_3$	$.2 * .2 * 1$	$* .9 * .8 * .4 * .6 = 0.006912$
		Total = 0.0096768


Evaluation Problem.

- **Evaluation problem.** Given the HMM $M=(A, B, \pi)$ and the observation sequence $O=O_1 O_2 \dots O_K$, calculate the probability that model M has generated sequence O .

- Trying to find probability of observations $O=O_1 O_2 \dots O_K$ by means of considering all hidden state sequences (as was done in example) is impractical:
 N^K hidden state sequences - exponential complexity.

- Use **Forward-Backward HMM algorithms** for efficient calculations.

- Define the forward variable $\alpha_k(i)$ as the joint probability of the partial observation sequence $O_1 O_2 \dots O_k$ and that the hidden state at time k is S_i : $\alpha_k(i) = P(O_1 O_2 \dots O_k, q_k = S_i)$



Forward recursion for HMM

- **Initialization:**
 $\alpha_k(i) = P(o_1, q_1 = s_i) = \pi_i b_i(o_1), 1 \leq i \leq N.$
- **Forward recursion:**

$$\alpha_{k+1}(i) = P(o_1, o_2, \dots, o_{k+1}, q_{k+1} = s_i) = \sum_j P(o_1, o_2, \dots, o_k, q_k = s_j, q_{k+1} = s_i) = \sum_j [\alpha_k(j) a_{ji}] b_i(o_{k+1}), 1 \leq j \leq N, 1 \leq k \leq K-1.$$
- **Termination:**
 $P(o_1, o_2, \dots, o_K) = \sum_i P(o_1, o_2, \dots, o_K, q_K = s_i) = \sum_i \alpha_K(i)$
- **Complexity:**
 $N^2 K$ operations.

Backward recursion for HMM

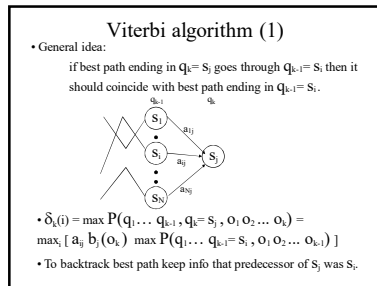
- Define the forward variable $\beta_k(i)$ as the joint probability of the partial observation sequence $o_{k+1} o_{k+2} \dots o_K$ given that the hidden state at time k is s_i : $\beta_k(i) = P(o_{k+1} o_{k+2} \dots o_K | q_k = s_i)$
- **Initialization:**
 $\beta_K(i) = 1, 1 \leq i \leq N.$
- **Backward recursion:**

$$\beta_k(i) = P(o_{k+1} o_{k+2} \dots o_K | q_k = s_i) = \sum_j P(o_{k+1} o_{k+2} \dots o_K, q_{k+1} = s_j | q_k = s_i) = \sum_j P(o_{k+2} o_{k+3} \dots o_K | q_{k+1} = s_j) a_{ji} b_j(o_{k+1}) = \sum_j \beta_{k+1}(j) a_{ji} b_j(o_{k+1}), 1 \leq j \leq N, 1 \leq k \leq K-1.$$
- **Termination:**
 $P(o_1, o_2, \dots, o_K) = \sum_i P(o_1, o_2, \dots, o_K, q_1 = s_i) = \sum_i P(o_1, o_2, \dots, o_K | q_1 = s_i) P(q_1 = s_i) = \sum_i \beta_1(i) b_i(o_1) \pi_i$

Decoding problem

- **Decoding problem.** Given the HMM $M=(A, B, \pi)$ and the observation sequence $O=o_1 o_2 \dots o_K$, calculate the most likely sequence of hidden states S_i that produced this observation sequence.
- We want to find the state sequence $Q=q_1 \dots q_K$ which maximizes $P(Q | o_1 o_2 \dots o_K)$, or equivalently $P(Q, o_1 o_2 \dots o_K)$.
- Brute force consideration of all paths takes exponential time. Use efficient **Viterbi algorithm** instead.
- Define variable $\delta_k(i)$ as the maximum probability of producing observation sequence $o_1 o_2 \dots o_k$ when moving along any hidden state sequence $q_1 \dots q_{k-1}$ and getting into $q_k = s_i$.

$$\delta_k(i) = \max P(q_1 \dots q_{k-1}, q_k = s_i, o_1 o_2 \dots o_k)$$
where max is taken over all possible paths $q_1 \dots q_{k-1}$.



Viterbi algorithm (2)

- **Initialization:**
 $\delta_1(i) = \max P(q_1 = s_i, o_1) = \pi_i b_i(o_1), 1 \leq i \leq N.$
- **Forward recursion:**

$$\delta_k(i) = \max P(q_1 \dots q_{k-1}, q_k = s_i, o_1 o_2 \dots o_k) = \max_j [a_{ji} b_i(o_k) \max P(q_1 \dots q_{k-1}, q_{k-1} = s_j, o_1 o_2 \dots o_{k-1})] = \max_j [a_{ji} b_i(o_k) \delta_{k-1}(j)], 1 \leq j \leq N, 2 \leq k \leq K.$$
- **Termination:** choose best path ending at time K
 $\max_i [\delta_K(i)]$
- Backtrack best path.

This algorithm is similar to the forward recursion of evaluation problem, with \sum replaced by max and additional backtracking.

Learning problem (1)

- **Learning problem.** Given some training observation sequences $O=o_1 o_2 \dots o_K$ and general structure of HMM (numbers of hidden and visible states), determine HMM parameters $M=(A, B, \pi)$ that best fit training data, that is maximizes $P(O|M)$.
- There is no algorithm producing optimal parameter values.
- Use iterative expectation-maximization algorithm to find local maximum of $P(O|M)$ - **Baum-Welch algorithm**.

Learning problem (2)

- If training data has information about sequence of hidden states (as in word recognition example), then use maximum likelihood estimation of parameters:

$$a_{ij} = P(s_j | s_i) = \frac{\text{Number of transitions from state } S_i \text{ to state } S_j}{\text{Number of transitions out of state } S_i}$$

$$b(v_m) = P(v_m | s_i) = \frac{\text{Number of times observation } V_m \text{ occurs in state } S_i}{\text{Number of times in state } S_i}$$

Baum-Welch algorithm

General idea:

$$a_{ij} = P(s_j | s_i) = \frac{\text{Expected number of transitions from state } S_i \text{ to state } S_j}{\text{Expected number of transitions out of state } S_i}$$

$$b(v_m) = P(v_m | s_i) = \frac{\text{Expected number of times observation } V_m \text{ occurs in state } S_i}{\text{Expected number of times in state } S_i}$$

$$\pi_i = P(s_i) = \text{Expected frequency in state } S_i \text{ at time } k=1.$$

Baum-Welch algorithm: expectation step(1)

- Define variable $\xi_k(i,j)$ as the probability of being in state S_i at time k and in state S_j at time $k+1$, given the observation sequence $O_1 O_2 \dots O_K$.

$$\xi_k(i,j) = P(q_k = S_i, q_{k+1} = S_j \mid O_1 O_2 \dots O_K)$$

$$\xi_k(i,j) = \frac{P(q_k = S_i, q_{k+1} = S_j, O_1 O_2 \dots O_K)}{P(O_1 O_2 \dots O_K)} = \frac{P(q_k = S_i, O_1 O_2 \dots O_k) a_j b_j(O_{k+1}) P(O_{k+2} \dots O_K \mid q_{k+1} = S_j)}{P(O_1 O_2 \dots O_K)} = \frac{\alpha_k(i) a_j b_j(O_{k+1}) \beta_{k+1}(j)}{\sum_i \alpha_k(i) a_i b_i(O_{k+1}) \beta_{k+1}(i)}$$

Baum-Welch algorithm: expectation step(2)

- Define variable $\gamma_k(i)$ as the probability of being in state S_i at time k , given the observation sequence $O_1 O_2 \dots O_K$.

$$\gamma_k(i) = P(q_k = S_i \mid O_1 O_2 \dots O_K)$$

$$\gamma_k(i) = \frac{P(q_k = S_i, O_1 O_2 \dots O_K)}{P(O_1 O_2 \dots O_K)} = \frac{\alpha_k(i) \beta_k(i)}{\sum_i \alpha_k(i) \beta_k(i)}$$

Baum-Welch algorithm: expectation step(3)

- We calculated $\xi_k(i,j) = P(q_k = S_i, q_{k+1} = S_j \mid O_1 O_2 \dots O_K)$ and $\gamma_k(i) = P(q_k = S_i \mid O_1 O_2 \dots O_K)$

- Expected number of transitions from state S_i to state $S_j = \sum_k \xi_k(i,j)$
- Expected number of transitions out of state $S_i = \sum_k \gamma_k(i)$

- Expected number of times observation V_m occurs in state $S_i = \sum_k \gamma_k(i)$, k is such that $O_k = V_m$

- Expected frequency in state S_i at time $k=1 : \gamma_1(i)$.

Baum-Welch algorithm: maximization step

$$a_j = \frac{\text{Expected number of transitions from state } S_j \text{ to state } S_i}{\text{Expected number of transitions out of state } S_j} = \frac{\sum_k \xi_k(i,j)}{\sum_k \gamma_k(i)}$$

$$b_i(V_m) = \frac{\text{Expected number of times observation } V_m \text{ occurs in state } S_i}{\text{Expected number of times in state } S_i} = \frac{\sum_k \xi_k(i,j)}{\sum_{k: O_k = V_m} \gamma_k(i)}$$

$$\pi_i = (\text{Expected frequency in state } S_i \text{ at time } k=1) = \gamma_1(i).$$